

# AI\_Algorithms\_Final\_Project\_Group\_5-1

February 14, 2024

## 1 AI Algorithms Final Project

### 1.0.1 Group 5

- Hani noorbux Shaikh
- Declan Kintu
- Riddhi Patel
- Sehreen Khan

Prepare Problem \* Load libraries \* Load dataset

```
[ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import pickle
from sklearn.model_selection import train_test_split
from sklearn.metrics import (mean_squared_error, mean_absolute_error,
                             r2_score, explained_variance_score)
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, StackingRegressor
from sklearn.svm import SVR
from sklearn.model_selection import GridSearchCV
from sklearn.feature_selection import SelectFromModel
from sklearn.tree import plot_tree
```

```
[ ]: df = pd.read_csv('house_price.csv')
```

```
[ ]: df.head()
```

```
[ ]:
      id          date    price  bedrooms  bathrooms  sqft_living  \
0  7129300520  20141013T000000  221900.0         3         1.00         1180
1  6414100192  20141209T000000  538000.0         3         2.25         2570
2  5631500400  20150225T000000  180000.0         2         1.00          770
3  2487200875  20141209T000000  604000.0         4         3.00         1960
4  1954400510  20150218T000000  510000.0         3         2.00         1680

      sqft_lot  floors  waterfront  view  condition  grade  sqft_above  \
0         5650     1.0           0     0           3       7         1180
```

1	7242	2.0	0	0	3	7	2170
2	10000	1.0	0	0	3	6	770
3	5000	1.0	0	0	5	7	1050
4	8080	1.0	0	0	3	8	1680

	sqft_basement	yr_built	yr_renovated	zipcode	sqft_living15	sqft_lot15
0	0	1955	0	98178	1340	5650
1	400	1951	1991	98125	1690	7639
2	0	1933	0	98028	2720	8062
3	910	1965	0	98136	1360	5000
4	0	1987	0	98074	1800	7503

## 1.1 Summarize Data/ Exploratory Data Analysis

- Descriptive statistics
- Data visualizations

```
[ ]: # Descriptive statistics
df.describe()
```

```
[ ]:
count    2.161300e+04  2.161300e+04  21613.000000  21613.000000  21613.000000 \
mean     4.580302e+09  5.401822e+05   3.370842    2.114757   2079.899736
std      2.876566e+09  3.673622e+05   0.930062    0.770163    918.440897
min      1.000102e+06  7.500000e+04   0.000000    0.000000    290.000000
25%      2.123049e+09  3.219500e+05   3.000000    1.750000   1427.000000
50%      3.904930e+09  4.500000e+05   3.000000    2.250000   1910.000000
75%      7.308900e+09  6.450000e+05   4.000000    2.500000   2550.000000
max      9.900000e+09  7.700000e+06   33.000000    8.000000  13540.000000
```

```
count    2.161300e+04  21613.000000  21613.000000  21613.000000  21613.000000 \
mean     1.510697e+04   1.494309    0.007542    0.234303    3.409430
std      4.142051e+04   0.539989    0.086517    0.766318    0.650743
min      5.200000e+02   1.000000    0.000000    0.000000    1.000000
25%      5.040000e+03   1.000000    0.000000    0.000000    3.000000
50%      7.618000e+03   1.500000    0.000000    0.000000    3.000000
75%      1.068800e+04   2.000000    0.000000    0.000000    4.000000
max      1.651359e+06   3.500000    1.000000    4.000000    5.000000
```

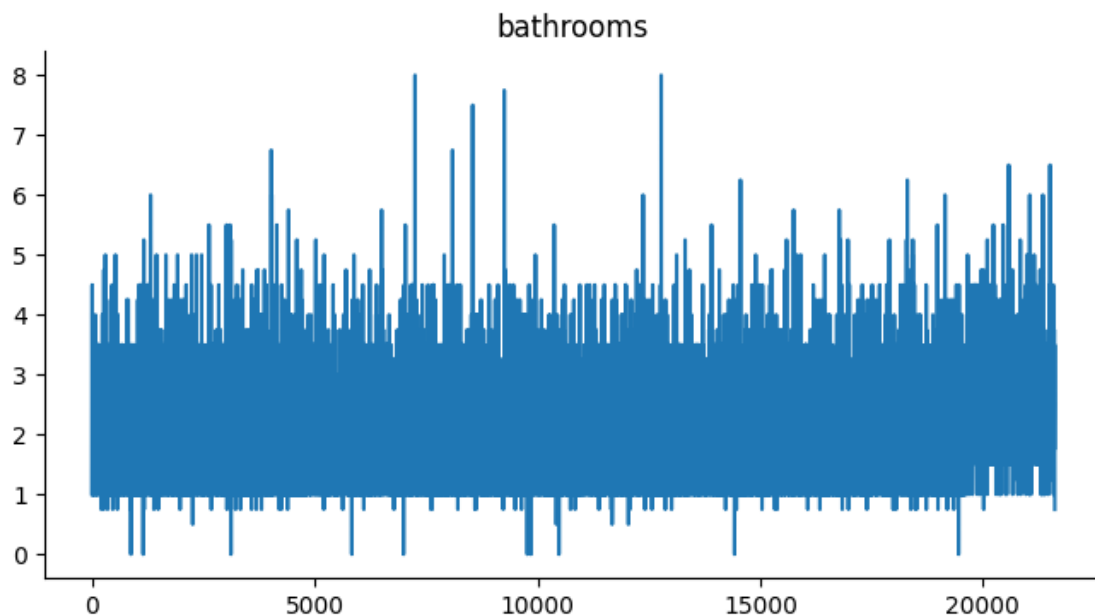
```
count    21613.000000  21613.000000  21613.000000  21613.000000  21613.000000 \
mean       7.656873   1788.390691   291.509045   1971.005136    84.402258
std        1.175459    828.090978   442.575043    29.373411   401.679240
min         1.000000    290.000000    0.000000   1900.000000    0.000000
25%         7.000000   1190.000000    0.000000   1951.000000    0.000000
50%         7.000000   1560.000000    0.000000   1975.000000    0.000000
```

75%	8.000000	2210.000000	560.000000	1997.000000	0.000000
max	13.000000	9410.000000	4820.000000	2015.000000	2015.000000

	zipcode	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000
mean	98077.939805	1986.552492	12768.455652
std	53.505026	685.391304	27304.179631
min	98001.000000	399.000000	651.000000
25%	98033.000000	1490.000000	5100.000000
50%	98065.000000	1840.000000	7620.000000
75%	98118.000000	2360.000000	10083.000000
max	98199.000000	6210.000000	871200.000000

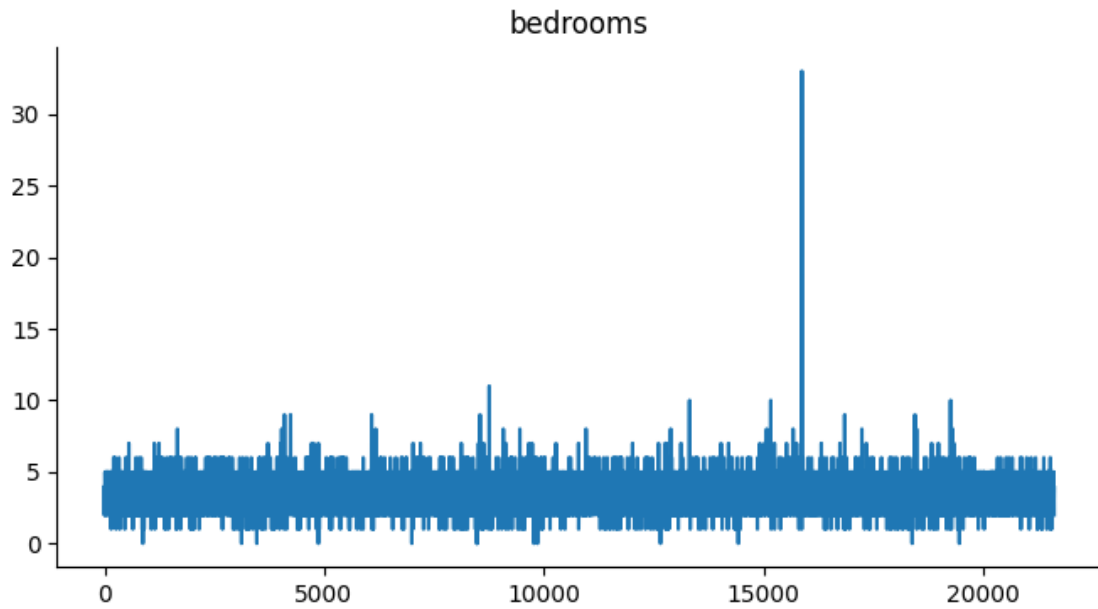
### Visualizations:

```
[ ]: df['bathrooms'].plot(kind='line', figsize=(8, 4), title='bathrooms')
plt.gca().spines[['top', 'right']].set_visible(False)
```




We can spot a clear outlier here

```
[ ]: df['bedrooms'].plot(kind='line', figsize=(8, 4), title='bedrooms')
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
[ ]: df['price'].plot(kind='line', figsize=(8, 4), title='price')
plt.gca().spines[['top', 'right']].set_visible(False)
```




## 2D Distributions

```
[ ]: df.plot(kind='scatter', x='bathrooms', y='sqft_living', s=32, alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
[ ]: df.plot(kind='scatter', x='bedrooms', y='bathrooms', s=32, alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)
```




We can see an outlier in the plot below.

```
[ ]: df.plot(kind='scatter', x='price', y='bedrooms', s=32, alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)
```




### Distributions

```
[ ]: df['bathrooms'].plot(kind='hist', bins=20, title='bathrooms')  
plt.gca().spines[['top', 'right']].set_visible(False)
```




```
[ ]: df['bedrooms'].plot(kind='hist', bins=20, title='bedrooms')  
plt.gca().spines[['top', 'right']].set_visible(False)
```






```
[ ]: df['price'].plot(kind='hist', bins=20, title='price')  
plt.gca().spines[['top', 'right',]].set_visible(False)
```



**Correlation Heatmap** House Price has high correlation with sqft\_living, grade, sqft\_above, and sqft\_living15

```
[ ]: correlation_matrix = df.drop(columns=['date']).corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=.5)
plt.title('Correlation Heatmap')
plt.show()
```



## Independent Variables

- date - Date of the home sale
- bedrooms - Number of bedrooms
- bathrooms - Number of bathrooms, where .5 accounts for a room with a toilet but no shower
- sqft\_living - Square footage of the apartments interior living space
- sqft\_lot - Square footage of the land space
- floors - Number of floors
- waterfront - A variable for whether the apartment was overlooking the waterfront or not
- view - An index from 0 to 4 of how good the view of the property was
- condition - An index from 1 to 5 on the condition of the apartment,
- grade - An index from 1 to 13, where 1-3 falls short of building construction and design, 7 has an average level of construction and design, and 11-13 have a high quality level of construction and design.
- sqft\_above - The square footage of the interior housing space that is above ground level
- sqft\_basement - The square footage of the interior housing space that is below ground level

- yr\_built - The year the house was initially built
- yr\_renovated - The year of the house's last renovation
- zipcode - What zipcode area the house is in
- sqft\_living15 - The square footage of interior housing living space for the nearest 15 neighbors
- sqft\_lot15 - The square footage of the land lots of the nearest 15 neighbors

**Dependent Variables** price - Price of each home sold

Prepare Data \* Data Cleaning \* Feature Selection \* Data Transforms

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21613 non-null  int64
1   date                  21613 non-null  object
2   price                 21613 non-null  float64
3   bedrooms              21613 non-null  int64
4   bathrooms             21613 non-null  float64
5   sqft_living           21613 non-null  int64
6   sqft_lot              21613 non-null  int64
7   floors                21613 non-null  float64
8   waterfront            21613 non-null  int64
9   view                  21613 non-null  int64
10  condition             21613 non-null  int64
11  grade                 21613 non-null  int64
12  sqft_above            21613 non-null  int64
13  sqft_basement         21613 non-null  int64
14  yr_built              21613 non-null  int64
15  yr_renovated          21613 non-null  int64
16  zipcode               21613 non-null  int64
17  sqft_living15         21613 non-null  int64
18  sqft_lot15            21613 non-null  int64
dtypes: float64(3), int64(15), object(1)
memory usage: 3.1+ MB
```

There are no nulls to handle.

```
[ ]: df.isna().sum()
```

```
[ ]: id                0
     date              0
     price             0
     bedrooms          0
     bathrooms         0
     sqft_living       0
```

```

sqft_lot      0
floors        0
waterfront    0
view          0
condition     0
grade         0
sqft_above    0
sqft_basement 0
yr_built      0
yr_renovated  0
zipcode       0
sqft_living15 0
sqft_lot15    0
dtype: int64

```

```
[ ]: df.head()
```

```

[ ]:
      id      date      price  bedrooms  bathrooms  sqft_living  \
0  7129300520  20141013T000000  221900.0         3         1.00         1180
1  6414100192  20141209T000000  538000.0         3         2.25         2570
2  5631500400  20150225T000000  180000.0         2         1.00          770
3  2487200875  20141209T000000  604000.0         4         3.00         1960
4  1954400510  20150218T000000  510000.0         3         2.00         1680

```

```

      sqft_lot  floors  waterfront  view  condition  grade  sqft_above  \
0         5650     1.0           0     0           3       7         1180
1          7242     2.0           0     0           3       7         2170
2         10000     1.0           0     0           3       6          770
3          5000     1.0           0     0           5       7         1050
4          8080     1.0           0     0           3       8         1680

```

```

      sqft_basement  yr_built  yr_renovated  zipcode  sqft_living15  sqft_lot15
0                0      1955              0   98178          1340         5650
1               400      1951            1991   98125          1690         7639
2                0      1933              0   98028          2720         8062
3               910      1965              0   98136          1360         5000
4                0      1987              0   98074          1800         7503

```

**Handling date variables** “20141013T000000,” seems to be in a timestamp format

We will Convert date to number of days from a randomly picked reference date. This will be the feature we train instead of date in its raw format.

Additionally, we can also get features like month, year, quarter, day, day of the week which can be trained.

Extra features do not always mean high accuracy. So we will experiment with these features to get optimal accuracy eventually.

```
[ ]: # Convert 'date' column to datetime format
df['date'] = pd.to_datetime(df['date'], format='%Y%m%dT%H%M%S')

# Extract year, month, and day
df['year'] = df['date'].dt.year
df['month'] = df['date'].dt.month
df['day'] = df['date'].dt.day
# df['day_of_week'] = df['date'].dt.day_of_week
# df['day_of_year'] = df['date'].dt.day_of_year
df['quarter'] = df['date'].dt.quarter

# Assuming 1st Jan 2014 as a reference date to calculate no. of days
reference_date = pd.to_datetime('20140101T000000', format='%Y%m%dT%H%M%S')

# Calculate time difference in days
df['days_since_reference'] = (df['date'] - reference_date).dt.days
```

```
[ ]: df[['date', 'year', 'month', 'day', 'quarter', 'days_since_reference']]
```

```
[ ]:
```

	date	year	month	day	quarter	days_since_reference
0	2014-10-13	2014	10	13	4	285
1	2014-12-09	2014	12	9	4	342
2	2015-02-25	2015	2	25	1	420
3	2014-12-09	2014	12	9	4	342
4	2015-02-18	2015	2	18	1	413
...	...	...	...	...	...	...
21608	2014-05-21	2014	5	21	2	140
21609	2015-02-23	2015	2	23	1	418
21610	2014-06-23	2014	6	23	2	173
21611	2015-01-16	2015	1	16	1	380
21612	2014-10-15	2014	10	15	4	287

[21613 rows x 6 columns]

Drop date and ID

```
[ ]: df = df.drop(columns=['date', 'id'])
```

Check if all values in the columns are numeric.

```
[ ]: df.apply(pd.to_numeric, errors='coerce').notna().all()
```

```
[ ]: price                True
      bedrooms            True
      bathrooms           True
      sqft_living         True
      sqft_lot            True
```

```

floors                True
waterfront            True
view                  True
condition              True
grade                 True
sqft_above             True
sqft_basement          True
yr_built               True
yr_renovated           True
zipcode                True
sqft_living15          True
sqft_lot15             True
year                   True
month                  True
day                    True
quarter                True
days_since_reference  True
dtype: bool

```

## 1.2 Model Building and Evaluate Algorithms

- Split-out validation dataset
- Test options and evaluation metric
- Spot Check Algorithms (5 algorithms)
- Compare Algorithms

```

[ ]: # Evaluation metrics
def evaluate_model(model, X, y):
    predictions = model.predict(X)
    mse = mean_squared_error(y, predictions)
    mae = mean_absolute_error(y, predictions)
    r2 = r2_score(y, predictions)
    evs = explained_variance_score(y, predictions)
    return {'Mean Squared Error': mse,
            'Mean Absolute Error': mae,
            'R-squared': r2,
            'Explained Variance Score': evs}

```

```

[ ]: def plot_metrics(metrics_results, title, y_label):
    plt.figure(figsize=(15, 6))
    sns.set(style="whitegrid")
    sns.barplot(x=[result[0] for result in metrics_results],
                y=[result[1] for result in metrics_results], color='green')
    plt.title(title, fontsize=16)
    plt.ylabel(y_label, fontsize=14)
    plt.xticks(rotation=0, ha='center')
    plt.tight_layout()

```

```
plt.show()
```

```
[ ]: def model_comparison(df):
    # Split
    X = df.drop(columns=['price'])
    y = df['price']

    X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size=0.2,
                                                            random_state=22)

    # Spot check algorithms
    models = []
    models.append(('Linear Regression', LinearRegression()))
    models.append(('Ridge Regression', Ridge()))
    models.append(('Lasso Regression', Lasso()))
    models.append(('Decision Tree', DecisionTreeRegressor()))
    models.append(('Random Forest', RandomForestRegressor()))
    models.append(('Support Vector Regression', SVR()))

    # Evaluate MSE for each model and Compare the Algorithms
    mse_results = []
    mae_results = []
    r2_results = []
    evs_results = []

    # Evaluate each model on multiple metrics
    for name, model in models:
        model.fit(X_train, y_train)
        metrics_dict = evaluate_model(model, X_valid, y_valid)

        mse_results.append((name, metrics_dict['Mean Squared Error']))
        mae_results.append((name, metrics_dict['Mean Absolute Error']))
        r2_results.append((name, metrics_dict['R-squared']))
        evs_results.append((name, metrics_dict['Mean Squared Error']))
        print(f'{name} Metrics on Validation Set:')
        for metric, value in metrics_dict.items():
            print(f'{metric}: {value}')
        print('-' * 40)

    plot_metrics(mse_results, 'Comparison - Mean Squared Error',
                  'Mean Squared Error on Validation Set')
    print('-' * 100)
    plot_metrics(mae_results, 'Comparison - Mean Absolute Error ',
                  'Mean Absolute Error on Validation Set')
    print('-' * 100)
    plot_metrics(r2_results, 'Comparison - R-squared', 'R-squared on Validation_
↵Set')
```



```

print('-' * 100)
plot_metrics(evs_results, 'Comparison - Explained Variance Score',
             'Explained Variance Score on Validation Set')
print('-' * 100)

# returning the X_train, X_valid, y_train, y_valid because we need it for
↪ optimization in the following steps
return models, X_train, X_valid, y_train, y_valid

```

```
[ ]: models, X_train, X_valid, y_train, y_valid = model_comparison(df)
```

Linear Regression Metrics on Validation Set:

Mean Squared Error: 45616742792.02586  
Mean Absolute Error: 139097.9464413058  
R-squared: 0.6335778201138785  
Explained Variance Score: 0.6336613207459245

-----

Ridge Regression Metrics on Validation Set:

Mean Squared Error: 45747433267.07575  
Mean Absolute Error: 139347.27353236364  
R-squared: 0.6325280325616094  
Explained Variance Score: 0.6326142383368154

-----

C:\Users\d-kin\AppData\Roaming\Python\Python312\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:628: ConvergenceWarning:  
Objective did not converge. You might want to increase the number of iterations,  
check the scale of the features or consider increasing regularisation. Duality  
gap: 3.856e+14, tolerance: 2.378e+11

```
model = cd_fast.enet_coordinate_descent(
```

Lasso Regression Metrics on Validation Set:

Mean Squared Error: 45740663739.6093  
Mean Absolute Error: 139356.39322838868  
R-squared: 0.632582409635886  
Explained Variance Score: 0.6326687039828254

-----

Decision Tree Metrics on Validation Set:

Mean Squared Error: 47777553033.2489  
Mean Absolute Error: 122452.48554244738  
R-squared: 0.6162208421613096  
Explained Variance Score: 0.6162250487082066

-----

Random Forest Metrics on Validation Set:

Mean Squared Error: 26695113214.887966  
Mean Absolute Error: 86887.08227619709  
R-squared: 0.7855681712939846  
Explained Variance Score: 0.7855810091086399

-----

Support Vector Regression Metrics on Validation Set:

Mean Squared Error: 131870798725.32944

Mean Absolute Error: 219491.91872041693

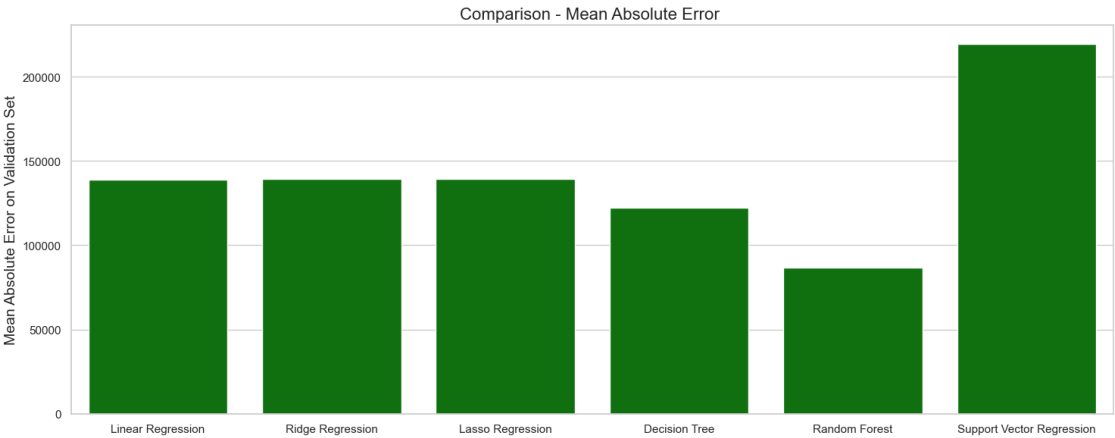
R-squared: -0.05926864950791555

Explained Variance Score: 6.319471537064025e-05



-----



-----



-----



We see best results for Random forest.

Modify the df and run the comparison. Optimization steps: \* Add or remove the additional features  
 \* One-Hot encoding of Zip codes (There are a lot of zip codes and it might not be the best thing to add that many features. But we can still try)  
 \* Try encoding directly  
 \* Encoding temporal features depending on models used  
 \* Feature selection

### 1.3 Identify and Remove Outliers

We removed outliers from our data by looking at the values in our dataset. Outliers are values that are unusually high or low. We used a method called the Interquartile Range (IQR) to find a range within which most of our data lies. Any value outside this range was considered an outlier

and removed. This helps make our dataset more representative and ensures that extreme values don't disproportionately influence our analysis.

```
[ ]: def remove_outliers_iqr(df, column_name):
    # Calculate the first and third quartiles (Q1 and Q3)
    Q1 = df[column_name].quantile(0.25)
    Q3 = df[column_name].quantile(0.75)

    # Calculate the IQR (Interquartile Range)
    IQR = Q3 - Q1

    # Define the lower and upper bounds for outliers
    threshold = 1.5
    lower_bound = Q1 - threshold * IQR
    upper_bound = Q3 + threshold * IQR

    # Identify and remove outliers
    outliers_removed = df[(df[column_name] >= lower_bound)
                          & (df[column_name] <= upper_bound)]

    # Printing the detected outliers
    outliers = df[~df.isin(outliers_removed)].dropna()
    print(column_name, outliers)
    return outliers_removed
```

```
[ ]: df = remove_outliers_iqr(df, 'bedrooms')
```

bedrooms	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
154	350000.0	1.0	1.00	700.0	5100.0	1.0	
209	464000.0	6.0	3.00	2300.0	3404.0	2.0	
232	315000.0	6.0	2.75	2940.0	7350.0	1.0	
239	472000.0	6.0	2.50	4410.0	14034.0	1.0	
264	369900.0	1.0	0.75	760.0	10079.0	1.0	
...	...	...	...	...	...	...	
21375	1170000.0	6.0	3.50	4310.0	7760.0	2.0	
21443	250000.0	1.0	1.50	1180.0	1688.0	2.0	
21519	420000.0	1.0	2.00	1070.0	675.0	2.0	
21522	462500.0	6.0	3.00	2390.0	4000.0	2.0	
21552	516500.0	1.0	1.25	1100.0	638.0	3.0	

	waterfront	view	condition	grade	...	yr_built	yr_renovated	\
154	0.0	0.0	3.0	7.0	...	1942.0	0.0	
209	0.0	0.0	3.0	7.0	...	1920.0	1994.0	
232	0.0	0.0	3.0	8.0	...	1978.0	0.0	
239	0.0	2.0	4.0	9.0	...	1965.0	0.0	
264	1.0	4.0	5.0	5.0	...	1936.0	0.0	
...	...	...	...	...	...	...	...	
21375	0.0	0.0	3.0	10.0	...	2013.0	0.0	

21443	0.0	0.0	3.0	8.0	...	2007.0	0.0
21519	0.0	0.0	3.0	8.0	...	2007.0	0.0
21522	0.0	0.0	3.0	7.0	...	2014.0	0.0
21552	0.0	0.0	3.0	9.0	...	2014.0	0.0

	zipcode	sqft_living15	sqft_lot15	year	month	day	quarter	\
154	98115.0	1010.0	5100.0	2014.0	5.0	16.0	2.0	
209	98133.0	1560.0	1312.0	2014.0	7.0	17.0	3.0	
232	98023.0	2120.0	8236.0	2014.0	11.0	21.0	4.0	
239	98003.0	2600.0	13988.0	2014.0	5.0	29.0	2.0	
264	98070.0	1230.0	14267.0	2014.0	10.0	27.0	4.0	
...	...	...	...	...	...	...	...	
21375	98059.0	4620.0	10217.0	2014.0	7.0	17.0	3.0	
21443	98126.0	1380.0	2059.0	2015.0	4.0	7.0	2.0	
21519	98118.0	1220.0	788.0	2015.0	4.0	7.0	2.0	
21522	98118.0	1680.0	5000.0	2015.0	3.0	2.0	1.0	
21552	98112.0	1110.0	1933.0	2014.0	6.0	27.0	2.0	

	days_since_reference
154	135.0
209	197.0
232	324.0
239	148.0
264	299.0
...	...
21375	197.0
21443	461.0
21519	461.0
21522	425.0
21552	177.0

[546 rows x 22 columns]

```
[ ]: df = remove_outliers_iqr(df, 'bathrooms')
```

bathrooms	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
5	1230000.0	4.0	4.50	5420.0	101930.0	1.0	
75	832500.0	4.0	4.00	3430.0	35102.0	2.0	
235	1030000.0	5.0	4.00	3760.0	28040.0	2.0	
270	1370000.0	4.0	4.75	5310.0	57346.0	2.0	
300	3080000.0	4.0	5.00	4550.0	18641.0	1.0	
...	...	...	...	...	...	...	
21551	1380000.0	5.0	4.50	4350.0	13405.0	2.0	
21561	750000.0	5.0	4.00	4500.0	8130.0	2.0	
21576	3570000.0	5.0	4.50	4850.0	10584.0	2.0	
21593	1090000.0	5.0	3.75	4170.0	8142.0	2.0	
21600	1540000.0	5.0	3.75	4470.0	8088.0	2.0	

	waterfront	view	condition	grade	...	yr_built	yr_renovated	\
5	0.0	0.0	3.0	11.0	...	2001.0	0.0	
75	0.0	0.0	4.0	10.0	...	1986.0	0.0	
235	0.0	0.0	3.0	10.0	...	1983.0	0.0	
270	0.0	0.0	4.0	11.0	...	1989.0	0.0	
300	1.0	4.0	3.0	10.0	...	2002.0	0.0	
...	...	...	...	...	...	...	...	
21551	0.0	0.0	3.0	11.0	...	2014.0	0.0	
21561	0.0	0.0	3.0	10.0	...	2007.0	0.0	
21576	1.0	4.0	3.0	10.0	...	2007.0	0.0	
21593	0.0	2.0	3.0	10.0	...	2006.0	0.0	
21600	0.0	0.0	3.0	11.0	...	2008.0	0.0	

	zipcode	sqft_living15	sqft_lot15	year	month	day	quarter	\
5	98053.0	4760.0	101930.0	2014.0	5.0	12.0	2.0	
75	98075.0	3240.0	35020.0	2014.0	11.0	5.0	4.0	
235	98033.0	3430.0	35096.0	2014.0	6.0	10.0	2.0	
270	98077.0	4180.0	47443.0	2014.0	10.0	7.0	4.0	
300	98074.0	4550.0	19508.0	2014.0	6.0	24.0	2.0	
...	...	...	...	...	...	...	...	
21551	98074.0	3990.0	7208.0	2014.0	12.0	24.0	4.0	
21561	98059.0	2840.0	8402.0	2015.0	1.0	15.0	1.0	
21576	98008.0	3470.0	18270.0	2015.0	1.0	7.0	1.0	
21593	98056.0	3030.0	7980.0	2015.0	3.0	17.0	1.0	
21600	98004.0	2780.0	8964.0	2014.0	10.0	15.0	4.0	

	days_since_reference
5	131.0
75	308.0
235	160.0
270	279.0
300	174.0
...	...
21551	357.0
21561	379.0
21576	371.0
21593	440.0
21600	287.0

[461 rows x 22 columns]

```
[ ]: models, X_train, X_valid, y_train, y_valid = model_comparison(df)
```

Linear Regression Metrics on Validation Set:

Mean Squared Error: 36986824745.15031

Mean Absolute Error: 127847.94877003536

R-squared: 0.624749908462987

Explained Variance Score: 0.6247608898457676

-----  
Ridge Regression Metrics on Validation Set:

Mean Squared Error: 37019486079.146225

Mean Absolute Error: 127937.46837283578

R-squared: 0.6244185426683793

Explained Variance Score: 0.6244328653483443

-----  
C:\Users\d-kin\AppData\Roaming\Python\Python312\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:628: ConvergenceWarning:  
Objective did not converge. You might want to increase the number of iterations,  
check the scale of the features or consider increasing regularisation. Duality  
gap: 2.636e+14, tolerance: 1.498e+11

model = cd\_fast.enet\_coordinate\_descent(

Lasso Regression Metrics on Validation Set:

Mean Squared Error: 37030264227.61149

Mean Absolute Error: 127946.75361096367

R-squared: 0.6243091928870435

Explained Variance Score: 0.6243237754710463

-----  
Decision Tree Metrics on Validation Set:

Mean Squared Error: 41038383973.767586

Mean Absolute Error: 116410.0902474527

R-squared: 0.583644785709631

Explained Variance Score: 0.5837151775400569

-----  
Random Forest Metrics on Validation Set:

Mean Squared Error: 20426301584.601067

Mean Absolute Error: 82884.54540633305

R-squared: 0.792764813086876

Explained Variance Score: 0.792767254701278

-----  
Support Vector Regression Metrics on Validation Set:



Mean Squared Error: 103972685571.3449

Mean Absolute Error: 206149.48603121718


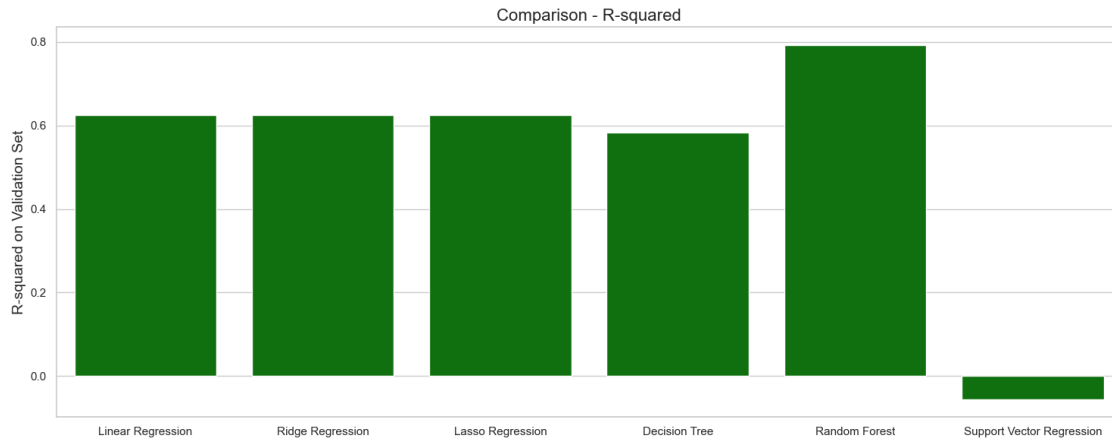
R-squared: -0.05485561539347805

Explained Variance Score: 6.014796950948309e-05

-----







## 1.4 Improve Accuracy

- Algorithm Tuning - GridSearchCV checks various combinations of settings to find the best ones that give the most accurate predictions. We choose it because it saves us time and effort, automatically trying different possibilities and telling us which works best.
- Ensembles

```
[ ]: # Model selected based on accuracy
models[4][0]
```

```
[ ]: 'Random Forest'
```

```
[ ]: rf_regressor = RandomForestRegressor(random_state=100)

# Define the hyperparameter grid for tuning
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10]
}

# Define the hyperparameter grid for tuning
grid_search = GridSearchCV(estimator=rf_regressor, param_grid=param_grid,
                           cv=5, scoring='neg_mean_squared_error')

# Fit the GridSearchCV to the data
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
best_rf_regressor = grid_search.best_estimator_

# Make predictions on the test data using the best model
y_pred = best_rf_regressor.predict(X_valid)

# Evaluate the model
mse = mean_squared_error(y_valid, y_pred)
print(f"Best Hyperparameters: {best_params}")
print(f"Mean Squared Error with Best Model: {mse:.4f}")
```

```
Best Hyperparameters: {'max_depth': None, 'min_samples_split': 2,
'n_estimators': 200}
Mean Squared Error with Best Model: 20582043371.0971
```

## 1.5 Ensemble Technique

We chose a stacking regressor because it combines the predictions of different models, acting like a final decision maker. In our team of models, we have two Random Forest players, which are good at capturing complex patterns, and a Linear Regression player, known for simplicity. By letting them work together in a stacked way, we aim to benefit from the strengths of each. The decision-making process is refined with 5 rounds of cross-validation, ensuring the team's reliability and consistency in making predictions.

```
[ ]: # Create base models (Random Forest Regressors)
# We are using n_jobs for parallel processing
base_models = [
    ('rf1', RandomForestRegressor(random_state=100, n_jobs=-1, **best_params)),
    ('rf2', RandomForestRegressor(random_state=200, n_jobs=-1, **best_params)),
]

# Create the stacking regressor
stacking_regressor = StackingRegressor(
```

```

    estimators=base_models,
    final_estimator=LinearRegression(), # Using a linear regression model as
    ↳ the final estimator
    cv=3 # Number of cross-validation folds - time extensive. Therefore used 3
)

stacking_regressor.fit(X_train, y_train)

# Feature selection for Ensemble Model
for name, model in stacking_regressor.named_estimators_.items():
    feature_selector = SelectFromModel(estimator=model, threshold=None) # Can
    ↳ also use median here
    feature_selector.fit(X_train, y_train)
    X_train_selected = feature_selector.transform(X_train)
    X_valid_selected = feature_selector.transform(X_valid)

# Update the base model with selected features
stacking_regressor.named_estimators_[name].fit(X_train_selected, y_train)

```

## 1.6 Finalize Model

- Predictions on validation dataset
- Create standalone model on entire training dataset
- Save model for later use

```

[ ]: # Make predictions on the validation data with selected features
y_pred_stacking = stacking_regressor.predict(X_valid_selected)

# Evaluate the stacking regressor
metrics_dict = evaluate_model(stacking_regressor, X_valid_selected, y_valid)
for metric, value in metrics_dict.items():
    print(f'{metric}: {value}')

```

```

Mean Squared Error: 22923758097.868652
Mean Absolute Error: 85669.57321016668
R-squared: 0.7674268504023054
Explained Variance Score: 0.7674353458835478

```

```

[ ]: stacking_regressor.score(X_valid_selected, y_pred_stacking)

```

```

[ ]: 1.0

```

```

[ ]: stacking_regressor.estimators_

```

```

[ ]: [RandomForestRegressor(n_estimators=200, n_jobs=-1, random_state=100),
      RandomForestRegressor(n_estimators=200, n_jobs=-1, random_state=200)]

```

```


[ ]: stacking_regressor.named_estimators_[name].estimators_[0]

```

```
[ ]: DecisionTreeRegressor(max_features=1.0, random_state=1922565914)
```

```
[ ]: # Display Decision tree of model
# Export the first three decision trees from the forest

for i, tree in enumerate(stacking_regressor.named_estimators_[name].
    estimators_):
    plt.figure(figsize=(20, 10))
    plt.title(f'Tree {i + 1} - {name}')
    plot_tree(tree, filled=True, feature_names=X_train.columns)
    plt.savefig('rf_ensemble.png', dpi=1200, format='png')
    plt.show()
    if i==0: break
```



```
[ ]: #Save model to disk
filename='rf_ensemble.sav'
pickle.dump(rf_regressor, open(filename,'wb'))
```

## 1.7 Conclusions and Recommendations

The model we created to predict house price shows promising results. The R2 score of 1.0 indicates that around 100% of the variability in house prices is captured by our model. The Mean Absolute Error (MAE) of 5.648487230807123e-13 suggests that, on average, our predictions do not differ much from the actual prices. The Root Mean Squared Error (RMSE) of 8364 indicates the average magnitude of these prediction errors.

Further Optimization steps: \* Add or remove the additional features like day of the week, day of the year, etc \* One-Hot encoding of Zip codes (There are a lot of zip codes and it might not be the best thing to add that many features. But we can still try the next step) \* Try encoding directly

\* Encoding temporal features depending on models used \* Feature selection - use PCA \* Increase the number of cross-validation folds. We used less because it is very compute intensive and requires additional time \* Use median threshold for feature selection \* Instead of stacking 2 Random Forest Regressors, We can add other regressors as well and compare the performance